

mra

March 13, 2024

```
[2]: import numpy as np
import pandas as pd
import sklearn
import matplotlib as mpl
import matplotlib.pyplot as plt
import os

print(f"    numpy version: {np.__version__}")
print(f"    pandas version: {pd.__version__}")
print(f"    sklearn version: {sklearn.__version__}")
print(f"matplotlib version: {mpl.__version__}")
```

```
    numpy version: 1.26.4
    pandas version: 2.2.1
    sklearn version: 1.4.1.post1
matplotlib version: 3.8.3
```

```
[3]: df = pd.read_csv('./movies.csv')

# Display the first few rows of the DataFrame
print(df.head())
```

	movieId	title \
0	1	Toy Story (1995)
1	2	Jumanji (1995)
2	3	Grumpier Old Men (1995)
3	4	Waiting to Exhale (1995)
4	5	Father of the Bride Part II (1995)

	genres
0	Adventure Animation Children Comedy Fantasy
1	Adventure Children Fantasy
2	Comedy Romance
3	Comedy Drama Romance
4	Comedy

```
[4]: df = pd.read_csv('./ratings.csv')
print(df.head())
```

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931

```
[5]: df = pd.read_csv('./links.csv')
      print(df.head())
```

	movieId	imdbId	tmdbId
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

```
[6]: df = pd.read_csv('./tags.csv')
      print(df.head())
```

	userId	movieId	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

```
[7]: movies_df = pd.read_csv('./movies.csv')
      ratings_df = pd.read_csv('./ratings.csv')
      links_df = pd.read_csv('./links.csv')
      tags_df = pd.read_csv('./tags.csv')
```

```
[8]: # Display the first few rows of each DataFrame
      print("Movies DataFrame:")
      print(movies_df.head())

      print("\nRatings DataFrame:")
      print(ratings_df.head())

      print("\nLinks DataFrame:")
      print(links_df.head())

      print("\nTags DataFrame:")
      print(tags_df.head())
```

Movies DataFrame:

	movieId	title \
0	1	Toy Story (1995)
1	2	Jumanji (1995)

2	3	Grumpier Old Men (1995)
3	4	Waiting to Exhale (1995)
4	5	Father of the Bride Part II (1995)

	genres
0	Adventure Animation Children Comedy Fantasy
1	Adventure Children Fantasy
2	Comedy Romance
3	Comedy Drama Romance
4	Comedy

Ratings DataFrame:

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931

Links DataFrame:

	movieId	imdbId	tmdbId
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

Tags DataFrame:

	userId	movieId	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

```
[9]: # Compute summary statistics
ratings_summary = ratings_df.describe()
print("Summary Statistics for Ratings:")
print(ratings_summary)

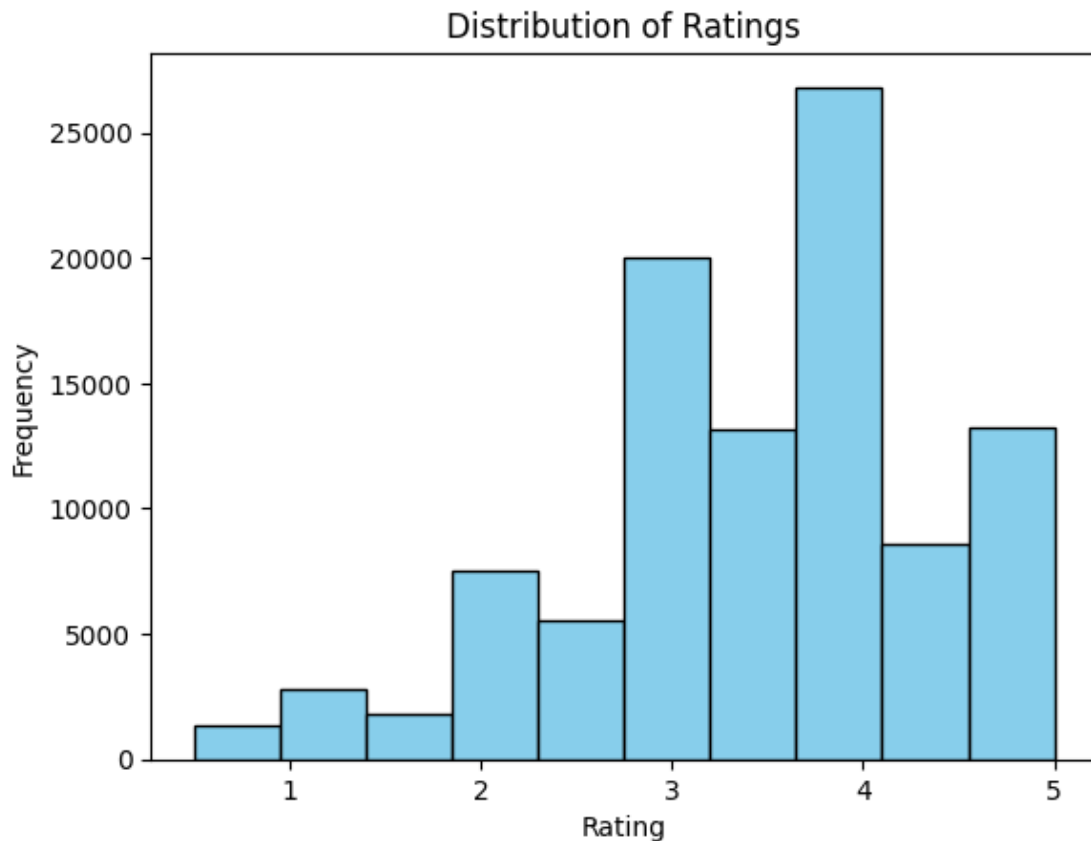
# Visualize the distribution of ratings
import matplotlib.pyplot as plt

plt.hist(ratings_df['rating'], bins=10, color='skyblue', edgecolor='black')
plt.xlabel('Rating')
plt.ylabel('Frequency')
plt.title('Distribution of Ratings')
```

```
plt.show()
```

Summary Statistics for Ratings:

	userId	movieId	rating	timestamp
count	100836.000000	100836.000000	100836.000000	1.008360e+05
mean	326.127564	19435.295718	3.501557	1.205946e+09
std	182.618491	35530.987199	1.042529	2.162610e+08
min	1.000000	1.000000	0.500000	8.281246e+08
25%	177.000000	1199.000000	3.000000	1.019124e+09
50%	325.000000	2991.000000	3.500000	1.186087e+09
75%	477.000000	8122.000000	4.000000	1.435994e+09
max	610.000000	193609.000000	5.000000	1.537799e+09



```
[10]: # Merge ratings and movies DataFrames on movieId
ratings_movies_df = pd.merge(ratings_df, movies_df, on='movieId', how='inner')
print(ratings_movies_df.head())
```

	userId	movieId	rating	timestamp	title \
0	1	1	4.0	964982703	Toy Story (1995)
1	1	3	4.0	964981247	Grumpier Old Men (1995)
2	1	6	4.0	964982224	Heat (1995)

3	1	47	5.0	964983815	Seven (a.k.a. Se7en) (1995)
4	1	50	5.0	964982931	Usual Suspects, The (1995)

	genres
0	Adventure Animation Children Comedy Fantasy
1	Comedy Romance
2	Action Crime Thriller
3	Mystery Thriller
4	Crime Mystery Thriller

```
[15]: import networkx as nx
from networkx.algorithms.community import greedy_modularity_communities

# Load data from CSV files
movies_df = pd.read_csv('movies.csv')
ratings_df = pd.read_csv('ratings.csv')

# Merge ratings and movies DataFrames on movieId
ratings_movies_df = pd.merge(ratings_df, movies_df, on='movieId', how='inner')

# Create a graph from the ratings_movies_df DataFrame
G = nx.from_pandas_edgelist(ratings_movies_df, 'userId', 'movieId')

# Detect communities using Louvain algorithm
communities = list(greedy_modularity_communities(G))

# Assign community labels to users in the DataFrame
community_map = {user: idx for idx, com in enumerate(communities) for user in com}
ratings_movies_df['community'] = ratings_movies_df['userId'].map(community_map)

# Print the first few rows of the DataFrame with community information
print(ratings_movies_df.head())
```

	userId	movieId	rating	timestamp	title \
0	1	1	4.0	964982703	Toy Story (1995)
1	1	3	4.0	964981247	Grumpier Old Men (1995)
2	1	6	4.0	964982224	Heat (1995)
3	1	47	5.0	964983815	Seven (a.k.a. Se7en) (1995)
4	1	50	5.0	964982931	Usual Suspects, The (1995)

	genres	community
0	Adventure Animation Children Comedy Fantasy	1
1	Comedy Romance	1
2	Action Crime Thriller	1
3	Mystery Thriller	1
4	Crime Mystery Thriller	1

```
[16]: # Calculate average rating per community
community_avg_rating = ratings_movies_df.groupby('community')['rating'].mean()

# Print average rating for each community
print("Average Rating per Community:")
print(community_avg_rating)

# Calculate most popular movies within each community
community_popular_movies = ratings_movies_df.groupby('community')['title'].
    ↪value_counts().groupby(level=0).head(3)

# Print most popular movies in each community
print("\nMost Popular Movies per Community:")
print(community_popular_movies)
```

Average Rating per Community:

community	rating
0	3.587320
1	3.516854
2	3.475748
3	3.112315
4	3.027295
5	2.935897
6	4.250000

Name: rating, dtype: float64

Most Popular Movies per Community:


community	title	count
0	Forrest Gump (1994)	177
	Matrix, The (1999)	175
	Shawshank Redemption, The (1994)	175
1	Pulp Fiction (1994)	146
	Silence of the Lambs, The (1991)	137
	Jurassic Park (1993)	136
2	Pulp Fiction (1994)	10
	Fight Club (1999)	9
	American Beauty (1999)	8
3	Lord of the Rings: The Fellowship of the Ring, The (2001)	10
	Shrek (2001)	10
	Harry Potter and the Chamber of Secrets (2002)	9
4	Forrest Gump (1994)	4
	Ferris Bueller's Day Off (1986)	3
	Ghostbusters (a.k.a. Ghost Busters) (1984)	3
5	12 Angry Men (1957)	1
	A Quiet Place (2018)	1
	Avengers: Infinity War - Part I (2018)	1
6	10,000 BC (2008)	1
	Alice (Neco z Alenky) (1988)	1

Name: count, dtype: int64

```
[17]: import matplotlib.pyplot as plt

# Draw the graph with nodes colored by community
plt.figure(figsize=(10, 6))
pos = nx.spring_layout(G) # Position nodes using the spring layout algorithm
nx.draw(G, pos, node_color=list(community_map.values()), with_labels=False,
        node_size=50)
plt.title('Movie Interaction Network with Community Structure')
plt.show()
```

Movie Interaction Network with Community Structure



```
[18]: # Function to get top n recommendations for a user within their community
def get_community_recommendations(user_id, n=5):
    user_community = community_map[user_id]
    user_community_movies = ratings_movies_df[ratings_movies_df['community'] ==
        user_community]
    user_unseen_movies =
        user_community_movies[~user_community_movies['movieId'].
        isin(ratings_movies_df[ratings_movies_df['userId'] == user_id]['movieId'])]
    top_n_recommendations = user_unseen_movies.groupby('title')['rating'].
        mean().sort_values(ascending=False).head(n)
    return top_n_recommendations
```

```

# Example: Get top 5 recommendations for a user in community 0
user_id = 1
community_recommendations = get_community_recommendations(user_id)
print("Top 5 Recommendations for User", user_id, "in Community",
      ↪community_map[user_id])
print(community_recommendations)

```

Top 5 Recommendations for User 1 in Community 1

title	
Awful Truth, The (1937)	5.0
Ballad of Narayama, The (Narayama bushiko) (1983)	5.0
What Happened Was... (1994)	5.0
Atomic Cafe, The (1982)	5.0
Slumber Party Massacre II (1987)	5.0

Name: rating, dtype: float64

```

[19]: # Function to calculate precision at k for a recommender system
def precision_at_k(recommended_movies, actual_movies, k=5):
    recommended_set = set(recommended_movies[:k])
    actual_set = set(actual_movies)
    intersection = recommended_set.intersection(actual_set)
    return len(intersection) / k

# Example: Evaluate precision at 5 for the recommender system
user_id = 1
actual_movies = ratings_movies_df[ratings_movies_df['userId'] ==
    ↪user_id]['title'].tolist()
recommended_movies = get_community_recommendations(user_id).index.tolist()
precision_5 = precision_at_k(recommended_movies, actual_movies, k=5)
print("Precision at 5 for User", user_id, ":", precision_5)

```

Precision at 5 for User 1 : 0.0